

Representation Multiplicity in Causal Forests

Yi Niu*

September 2026

Abstract

Including covariates alongside strictly monotone transformations can change a causal forest’s treatment decisions without adding information. Random feature selection favors covariates represented by multiple columns. Under stated conditions, I show that this imbalance can persist as samples grow. Treatment effect components associated with other covariates are omitted, attenuated, or recovered depending on their inclusion probabilities and tree depth. Simulations and a job-training replication illustrate sensitivity to redundant variables. Sampling groups of variables that generate identical splits makes forest predictions invariant across equivalent representations on the data used to define the groups, with canonical routing and common randomness.

Keywords: causal forests, feature subsampling, policy learning, representation invariance, redundant covariates.

JEL Classification: C14, C21, C55

*Department of Economics, University of Pennsylvania. Email: n1y1@sas.upenn.edu

1 Introduction

Suppose two researchers use a causal forest to estimate how a job-training program affects later earnings as a function of baseline household income and age. One includes income and age. The other keeps age but includes income, log income, and several other strictly increasing recodings of income. If baseline income has positive support, these income columns generate the same threshold partitions and contain the same covariate information. A feature-subsampled forest nevertheless treats them as separate candidates, making an income split more likely to enter the node-level search and an age split less likely.

The effect of this change depends on the whole path to a leaf. A covariate that drives treatment-effect heterogeneity matters only after a split separates its relevant states. If such a covariate is represented once among many columns, each node offers a small chance to use it. Those chances accumulate with path depth.

The main result models two independent covariates with binary treatment-effect components. One has a smaller coefficient than the other. The first representation repeats the covariate with the smaller coefficient and leaves the other as a singleton. The second does the reverse. Both representations generate the same sigma-field and the same split-action library. Let H_n be effective path depth, M_n the number of columns, and Q_{M_n} the number sampled at a node. The access-depth index is $\mathcal{I}_n = H_n \mathbb{E}[Q_{M_n}]/M_n$.

Under the continuation and score-transfer conditions stated below, if $\mathcal{I}_n \rightarrow 0$, the singleton covariate remains unused along a prediction path with probability approaching one. Each forest then retains only its repeated component. The two conditional average treatment effect (CATE) limits have opposite signs on half the covariate population, and repeating the covariate with the smaller coefficient produces nonvanishing policy regret. If $\mathcal{I}_n \rightarrow \kappa \in (0, \infty)$, the singleton component enters with a strictly intermediate weight under the additional path-depth condition in Theorem 1. If $\mathcal{I}_n \rightarrow \infty$, both forests recover the CATE. Under these conditions, recovery holds for every fixed finite representation but not uniformly over growing, information-equivalent representations.

The candidate-count rule translates the access-depth index into a condition on how quickly the number of columns may grow. In the **grf** R package for generalized random forests, the **mtry** parameter controls how many covariate columns are randomly sampled as split candidates at each node. With **mtry** fixed, $\mathcal{I}_n$ has order H_n/M_n . Under the package-default square-root scaling in **grf**, it has order $H_n/\sqrt{M_n}$. The critical representation size is thus of order H_n in the first case and H_n^2 in the second.

With 500 baseline columns and eight equivalent recodings, the Monte Carlo forests disagree in CATE sign for 45.7 percent of evaluation observations. Changing only the for-

est seed gives 0.09 percent. The random-forest outcome-regression specifications in [Chernozhukov et al. \[2022\]](#) retain age, education, and two prior-earnings variables together with their squares. The replication removes split-equivalent columns from the outcome regression while holding the samples, folds, Riesz stage based on the least absolute shrinkage and selection operator (Lasso), and orthogonal score fixed.

The remedy is to sample split-action classes rather than raw column labels. On a declared finite array, columns with the same weak ordering and ties generate the same thresholded actions for classification and regression trees (CART). Collapsing them before candidate sampling gives identical predictions under canonical routing and common class-level randomness. In the Monte Carlo design, class sampling eliminates representation-induced sign disagreement and reduces policy regret, although it increases CATE mean squared error. In the application of [Chernozhukov et al. \[2022\]](#), grouping each raw variable with its split-equivalent square removes the extra candidate weight created by retaining both columns while preserving every certified CART threshold partition.

The paper builds on the honest causal-forest framework of [Wager and Athey \[2018\]](#) and the local-moment formulation of generalized random forests in [Athey et al. \[2019\]](#). These papers establish conditions for consistent estimation and valid inference using forest-based adaptive neighborhoods. I retain honest sample separation and examine how the covariate representation changes those neighborhoods. Honest sample separation does not by itself ensure that every treatment-effect component is recovered, because feature subsampling can keep a relevant covariate out of the split search.

[Louppe \[2014\]](#) explains how retaining redundant predictors changes their probability of entering the split search and can affect prediction. Here I study how that change in split access affects limiting CATE estimates and treatment-rule regret. [Chi et al. \[2022\]](#) and [Klusowski and Tian \[2024\]](#) provide access-dependent recovery bounds. [Mei et al. \[2026\]](#) also characterize bias from unresolved components for forests with training-data-independent partitions. I derive representation-specific prediction limits for sample-grown honest trees under continuation and score-transfer conditions. Their path lengths depend on earlier splits, so singleton access must be analyzed jointly with stopping. An additional failure-path depth condition identifies the partial-recovery weight.

[Chen and Zhang \[2013\]](#) sample modules before representative variables, while [Colot et al. \[2021\]](#) weight features inversely by cluster size. Here, groups are certified to generate identical threshold partitions on a declared array. Proposition 1 establishes exact invariance on this array under canonical routing and common class-level randomness.

Section 2 defines the model. Section 3 derives the forest limits and their treatment-rule consequences. Section 4 gives the class-sampling guarantee, and Section 5 reports Monte

Carlo results and a replication of the empirical analyses in [Chernozhukov et al. \[2022\]](#). All proofs are in the appendix.

2 Model Setup

Let $X = (X_1, X_2) \sim \text{Unif}([0, 1]^2)$ be the covariate vector. For coordinate $j \in \{1, 2\}$, define the binary state indicator

$$s_j(X) = 2\mathbf{1}\{X_j \geq 1/2\} - 1,$$

and define the CATE

$$\tau_a(X) = a_1 s_1(X) + a_2 s_2(X),$$

indexed by the coefficient vector $a = (a_1, a_2)$. Restrict a to the parameter class

$$\mathcal{C}_a = \{(a_1, a_2) : \underline{a} \leq a_1 < a_2 \leq \bar{a}, a_2 - a_1 \geq \Delta_a\}$$

where $\underline{a}$ and $\bar{a}$ are fixed lower and upper coefficient bounds satisfying $0 < \underline{a} < \bar{a}$, and Δ_a is a fixed minimum coefficient gap, with $0 < \Delta_a < \bar{a} - \underline{a}$. These restrictions keep the sign and regret comparisons nondegenerate uniformly over $\mathcal{C}_a$.

Let n be the sample size, and let $(Y_i, W_i, X_i)_{i=1}^n$ denote the i.i.d. observations. For treatment state $w \in \{0, 1\}$, the potential outcome $Y(w)$ follows

$$Y(w) = m_0(X) + w\tau_a(X) + U, \quad \mathbb{E}(U \mid X, W) = 0, \quad (1)$$

where m_0 is the baseline outcome function and U is the outcome disturbance. The binary treatment indicator W is randomized with propensity score one half, and the observed outcome Y satisfies

$$W \perp \{Y(0), Y(1)\} \mid X, \quad \mathbb{P}(W = 1 \mid X) = \frac{1}{2}, \quad Y = Y(W).$$

The baseline function m_0 and the conditional law of U are held fixed over $a \in \mathcal{C}_a$. Define the observed-outcome regression $m_a(X) = \mathbb{E}(Y \mid X) = m_0(X) + \tau_a(X)/2$. The oracle transformed-outcome score for observation i is

$$\Gamma_i = 4(W_i - 1/2)\{Y_i - m_a(X_i)\} \quad (2)$$

and satisfies $\mathbb{E}(\Gamma_i \mid X_i = x) = \tau_a(x)$ at every covariate value x . The fitted score $\hat{\Gamma}_i$ replaces m_a with the cross-fitted outcome-regression estimator $\hat{m}$.

For covariate index $j \in \{1, 2\}$, let $c_n \geq 1$ be the representation multiplicity. Choose strictly increasing recoding maps $\{\varphi_{j\ell n} : \ell = 1, \dots, c_n\}$, including the identity, and denote the ℓ -th recoded column by $X_j^{(\ell)} = \varphi_{j\ell n}(X_j)$. Define the paired covariate representations

$$R_{1,n}(X) = \{X_1^{(1)}, \dots, X_1^{(c_n)}, X_2\}, \quad R_{2,n}(X) = \{X_1, X_2^{(1)}, \dots, X_2^{(c_n)}\}.$$

Each representation has $M_n = c_n + 1$ raw columns. Both representations generate the covariate sigma-field $\sigma(X_1, X_2)$. A cut $X_j \leq t$ corresponds to $\varphi_{j\ell n}(X_j) \leq \varphi_{j\ell n}(t)$ in each recoded column. These corresponding cuts preserve partitions and routing in exact arithmetic.

Let s_n be the tree subsample size. Each tree draws s_n observations uniformly without replacement and randomly divides them into splitting and estimation samples, independently of their values. Both sample fractions are at least a fixed $\eta \in (0, 1/2)$, the minimum sample-allocation fraction. Write n_C for the splitting-sample count in node C . At a node with M raw columns, the tree draws candidate count $Q_M \in \{1, \dots, M\}$ and samples that many columns uniformly without replacement, independently of the data and previous candidate draws. It evaluates every distinct threshold partition of the sampled columns. With minimum child splitting-sample size k_n and split-balance parameter $\beta \in (0, 1/2)$, a split is admissible when each child contains at least $\max\{k_n, \beta n_C\}$ splitting observations. The tree chooses a maximum-gain admissible split, using a representation-independent order for exact ties. It stops if no sampled admissible action has positive gain, with no separate depth cap or stopping rule. The estimation sample supplies the mean fitted score in each leaf, with value zero assigned to an empty estimation leaf. Conditional on the training data, fitted nuisance functions, and root sample assignments, each node's action depends only on its ancestral states, its observations, and fresh randomness assigned to that node. Randomization is independent across node addresses.

For a generic tree under representation h , let $D_{h,n}(x)$ indicate that the singleton X_{3-h} has been used as a splitting coordinate along the path to x . Even after a split on X_{3-h} , the terminal leaf may contain both values of s_{3-h} . The proof controls the resulting approximation error in $L^1(P_X)$, where P_X is the covariate distribution and $\|f\|_{L^1(P_X)} = \int |f(x)| dP_X(x)$ is the integrated absolute norm. For a candidate split indexed by column j and threshold t , let C be the parent cell and C_L and C_R its left and right child cells. Define the pseudo-outcome split-gain criterion

$$\widehat{\mathcal{Q}}_C^\Gamma(j, t) = \widehat{p}_{L|C} \widehat{p}_{R|C} \left(\widehat{\Gamma}_{C_L} - \widehat{\Gamma}_{C_R} \right)^2. \quad (3)$$

Here $\widehat{p}_{L|C}$ and $\widehat{p}_{R|C}$ are the empirical child shares within C , and $\widehat{\Gamma}_{C_L}$ and $\widehat{\Gamma}_{C_R}$ are the corresponding child-cell score means. Let $\widehat{\mathcal{Q}}_C^L(j, t)$ be the forest algorithm's normalized split-gain criterion, $\bar{T}_{h,n}$ the arithmetic mean of the B_n honest score-leaf predictions, where B_n is

the number of trees, and $\widehat{T}_{h,n}$ the reported forest predictor under representation $R_{h,n}$. The following conditions connect the forest algorithm to the score and control tree depth.

Assumption 1. *All probability statements and rates hold uniformly over $a \in \mathcal{C}_a$. Outcomes are uniformly bounded, the conditional law of U has a continuous density on bounded support, and $\|\widehat{m} - m_a\|_\infty = o_p(1)$.*

For fixed $C_0 > 4$, with probability tending to one, every column admits a positive-gain admissible split at every reached node with $n_C \geq C_0 k_n$, on both ordinary and locally conditioned failure trees whenever defined. For fixed $\bar{s} < 1$,

$$s_n/n \leq \bar{s}, \quad k_n \rightarrow \infty, \quad k_n/s_n \rightarrow 0, \quad \frac{\log(B_n s_n)}{k_n} \rightarrow 0, \quad B_n \rightarrow \infty. \quad (4)$$

On ordinary and failure trees,

$$\sup_{C,j,t} \left| \widehat{Q}_C^L(j,t) - \widehat{Q}_C^r(j,t) \right| = o_p(1), \quad \max_{h=1,2} \|\widehat{T}_{h,n} - \check{T}_{h,n}\|_{L^1(P_X)} = o_p(1). \quad (5)$$

Conditional on the training data and fitted nuisance functions, trees have identical marginal distributions and form independent groups of uniformly bounded size.

Assumption 1 connects the algorithm's split choices and predictions to the treatment effect. The first condition in (5) ensures that the algorithm prefers a split with a meaningfully larger gain in the fitted treatment-effect score. The second allows its reported predictions to be analyzed through averages of that score within honest leaves. The continuation condition ensures that a tree keeps splitting while its nodes remain sufficiently large, including when the candidate set contains only a previously used covariate. Continued growth gives the other covariate further chances to enter the split search. Together, these conditions allow the theorem to relate tree depth and candidate inclusion probabilities to the treatment effect components that the forest recovers.

Finally, define the effective path depth H_n , the singleton's per-node inclusion probability ω_n , and the access-depth index $\mathcal{I}_n$ by

$$H_n = \log(s_n/k_n), \quad \omega_n = \frac{\mathbb{E}Q_{M_n}}{M_n}, \quad \mathcal{I}_n = H_n \omega_n. \quad (6)$$

3 Forest Limits and Treatment-Rule Distortion

For $\omega_n < 1$, define a failure tree by conditioning the ordinary transition at each visited node locally on not choosing the singleton as the splitting coordinate. A conditioned transition

can split on the repeated coordinate or stop. Its conditioning probability is at least $1 - \omega_n > 0$. Let $J_{h,n}^0(X)$ count its candidate-set draws at size-eligible nodes, including a final unsuccessful search if one occurs. In the intermediate regime $\mathcal{I}_n \rightarrow \kappa \in (0, \infty)$, $\omega_n \rightarrow 0$, so this construction is eventually defined. For each fixed $a \in \mathcal{C}_a$, additionally suppose

$$\frac{J_{h,n}^0(X)}{H_n} \rightarrow_p d_h(a) \in (0, \infty), \quad h = 1, 2, \quad (7)$$

where $d_h(a)$ is the limiting normalized failure-path depth under representation h . The convergence accounts for randomness in the training sample, the locally conditioned failure-tree construction, and an independent test point $X \sim P_X$.

Theorem 1 (Representation-dependent forest limits). *Under the model and forest construction in Section 2 and Assumption 1, the two endpoint results hold uniformly over $a \in \mathcal{C}_a$. If $\mathcal{I}_n \rightarrow 0$, then*

$$\max_{h=1,2} \|\widehat{T}_{h,n} - a_h s_h\|_{L^1(P_X)} \rightarrow_p 0. \quad (8)$$

If $\mathcal{I}_n \rightarrow \infty$, then

$$\max_{h=1,2} \|\widehat{T}_{h,n} - \tau_a\|_{L^1(P_X)} \rightarrow_p 0. \quad (9)$$

For fixed $a \in \mathcal{C}_a$, if $\mathcal{I}_n \rightarrow \kappa \in (0, \infty)$ and (7) holds, then

$$\begin{aligned} \widehat{T}_{1,n} &\rightarrow_p a_1 s_1 + \{1 - e^{-\kappa d_1(a)}\} a_2 s_2, \\ \widehat{T}_{2,n} &\rightarrow_p \{1 - e^{-\kappa d_2(a)}\} a_1 s_1 + a_2 s_2 \end{aligned} \quad \text{in } L^1(P_X). \quad (10)$$

When $\mathcal{I}_n \rightarrow 0$, the singleton remains unused with probability approaching one, and its component drops out of the limit. In the intermediate regime, $e^{-\kappa d_h(a)}$ is the limiting probability that a path reaches its leaf without using the singleton, so the singleton component is attenuated by the complementary factor $1 - e^{-\kappa d_h(a)} \in (0, 1)$. When the access-depth index diverges, both components are recovered in the forest prediction. The result applies equally to exact copies and to growing sequences of strictly increasing recodings.

For a binary treatment rule π , where $\pi(x) = 1$ assigns treatment and $\pi(x) = 0$ withholds it, define its policy value $V_a(\pi)$ and policy regret $\text{Reg}_a(\pi)$ by

$$V_a(\pi) = \mathbb{E}\{\tau_a(X)\pi(X)\}, \quad \text{Reg}_a(\pi) = V_a(\pi_a^*) - V_a(\pi).$$

Policy value is the expected treatment gain relative to treating no one, and policy regret is the value lost relative to the optimal unconstrained rule π_a^* . This optimal rule and the fitted

rule $\widehat{\pi}_{h,n}$ under representation h are

$$\pi_a^*(x) = \mathbf{1}\{\tau_a(x) > 0\} = \mathbf{1}\{s_2(x) = 1\}, \quad \widehat{\pi}_{h,n}(x) = \mathbf{1}\{\widehat{T}_{h,n}(x) > 0\}.$$

Write $\mathbb{P}_X$ for probability over an independent covariate draw, conditional on the fitted forests.

Corollary 1 (Sign reversal and policy regret). *Under the conditions of Theorem 1, if $\mathcal{I}_n \rightarrow 0$, then*

$$\mathbb{P}_X \left\{ \widehat{T}_{1,n}(X)\widehat{T}_{2,n}(X) < 0, \min_{h=1,2} |\widehat{T}_{h,n}(X)| \geq \underline{a}/2 \right\} \rightarrow_p \frac{1}{2}. \quad (11)$$

Moreover,

$$\text{Reg}_a(\widehat{\pi}_{2,n}) \rightarrow_p 0, \quad \text{Reg}_a(\widehat{\pi}_{1,n}) \rightarrow_p \frac{a_2 - a_1}{2}. \quad (12)$$

Thus the representation that repeats the covariate with the larger coefficient yields the optimal unconstrained rule in the selection regime. Repeating the covariate with the smaller coefficient yields a persistent treatment mistake.

The `grf` 2.4.0 algorithm reference [GRF Development Team, 2026] specifies the candidate-count law as

$$N_M \sim \text{Poisson}(\xi_M), \quad Q_M = \max\{1, \min(N_M, M)\}, \quad \xi_M = \text{mtry} \quad (13)$$

Here N_M is the uncapped Poisson candidate count, ξ_M is its mean parameter, and Q_M is the candidate count after truncation to $\{1, \dots, M\}$.

Corollary 2 (Representation-growth rates). *Under the conditions of Theorem 1, if $M_n \rightarrow \infty$ and $\xi_{M_n} \equiv \xi > 0$, then*

$$\omega_n \sim \frac{\xi + e^{-\xi}}{M_n},$$

so the critical representation size is $M_n \asymp H_n$. Under the package default $\xi_M = \min\{\lceil \sqrt{M} \rceil + 20, M\}$, $\omega_n \sim M_n^{-1/2}$, and the critical representation size is $M_n \asymp H_n^2$. Every fixed finite representation satisfies $\mathcal{I}_n \rightarrow \infty$.

Under the default rule, recovery is guaranteed when $M_n = o(H_n^2)$, compared with $M_n = o(H_n)$ for a fixed Poisson mean. Since $H_n = \log(s_n/k_n)$, increasing the sample size expands the scope for redundant columns only slowly. The default rule permits faster growth in the number of columns, but retaining many equivalent transformations can still leave a covariate represented once with few opportunities to affect prediction. Every fixed representation covered by the theorem eventually satisfies these growth conditions. In finite samples, however, these asymptotic rates provide no numerical cutoff for how many redundant columns are acceptable.

4 Finite-Array Class Sampling

I propose sampling split classes instead of individual columns. Fix an array $\mathcal{X}_{\text{cert}} = \{x_1, \dots, x_N\}$ containing every point used for fitting, prediction, variance estimation, and policy ranking. A column’s certificate is its weak ordering of this array, identified with the reversed ordering. Equal certificates generate the same unordered threshold partitions on every subset of the array.

Construct the classes before growing the forest. Replace each raw column on $\mathcal{X}_{\text{cert}}$ by its weak-rank vector, assigning the same rank to tied observations, and treat a rank vector and its reversal as the same certificate. Group columns with identical certificates, choose a canonical orientation and weak-rank representative for each group, and order the groups by their certificates. At each node, sample these groups uniformly without replacement and evaluate the usual CART threshold splits using their canonical representatives. Index child orientation, tie-breaking, folds, and all other randomization by the certificate rather than by the original column label. Hold nuisance fits fixed across representations or fit them on the same canonical matrix. All tuning rules and reported outputs use this matrix and the same settings. For policy ranking, break prediction ties by a fixed ordering of observation indices. Because this reduction occurs before candidate sampling, adding a recoding to an existing class, or removing one without removing its class, does not alter either the candidate probabilities or the random draws.

Proposition 1 (Class-sampled forest). *Fix the certified classes, canonical routing, and class-level candidate law. Under common class-indexed randomness, permuting columns, adding a column to an existing class, deleting a column without removing its class, complementing a column, or applying a recoding that preserves or reverses the observed weak ordering and ties leaves all tree partitions, forest weights, predictions, reported standard errors, and selected sets on $\mathcal{X}_{\text{cert}}$ unchanged.*

This guarantee applies to the declared array, not to future points whose ordering has not been certified. The grouping criterion also differs from that of [Chen and Zhang \[2013\]](#), who construct modules of correlated variables and sample a representative from each selected module. Correlation alone is not sufficient for the split classes used here.

5 Finite-Sample Evidence

The Monte Carlo compares CATE estimation error and policy regret across redundant encodings and under class sampling. In the job-training application of [Chernozhukov et al.](#)

[2022], the comparison instead concerns changes in individual outcome-regression contrasts and estimates of the average treatment effect on the treated (ATET). In each exercise, refitting with a different forest seed provides a benchmark for the changes caused by recoding.

5.1 Monte Carlo

Let p be the number of baseline covariates. The Monte Carlo independently draws $X \sim \text{Unif}([0, 1]^p)$, $W \sim \text{Bernoulli}(1/2)$, and

$$\begin{aligned}\tau(X) &= 0.20 + 0.65s_1 - 0.55s_2 + 0.10(X_3 - 0.5), \\ Y &= 1 + X_4 + 0.5X_5^2 - 0.5X_6 + W\tau(X) + \varepsilon, \quad \varepsilon \sim N(0, 1).\end{aligned}$$

Here ε is the outcome disturbance and $s_j = 2\mathbf{1}\{X_j \geq 1/2\} - 1$ is the binary state indicator for coordinate $j \in \{1, 2\}$. Table 1 compares paired representations that add eight monotone or order-reversing recodings of X_1 or X_2 . These are standardization, complementation, an affine change of units, a shifted logarithm, ranks, a shifted square, an exponential, and a shifted cube. Each cell averages 100 independent samples of 3,000 observations, with 30,000 evaluation observations and 2,000-tree forests fitted with `grf` 2.4.0. Both representations use the true nuisance functions $\mathbb{E}(Y \mid X)$ and $\mathbb{P}(W = 1 \mid X) = 1/2$. Thus the comparison isolates the forest fit from nuisance-estimation error. The forests use default `mtry`, a minimum node size of five, and honest half-samples with an equal split between partition construction and leaf estimation. Data and forest seeds are paired. For predictions $\hat{\tau}_1$ and $\hat{\tau}_2$, define the sign-disagreement rate SR_c as the fraction with opposite signs and both magnitudes at least the magnitude threshold $c \geq 0$.

Table 1: CATE sign disagreement under equivalent representations

Comparison	Raw columns	SR ₀ (%)	SR _{0.10} (%)
$p = 40$, representation change	48	8.81 (0.66)	0.02 (0.003)
$p = 100$, representation change	108	23.41 (0.66)	12.70 (0.83)
$p = 250$, representation change	258	38.60 (0.84)	26.87 (0.73)
$p = 500$, representation change	508	45.74 (0.54)	32.88 (0.91)
$p = 500$, forest-seed change	508	0.09 (0.01)	0.00 (0.00)
$p = 500$, class-sampled forest	508	0.00 (0.00)	0.00 (0.00)

Notes: Entries are means across 100 independent replications, with Monte Carlo standard errors in parentheses. The first four rows compare representations with eight added encodings. The last two rows change only the seed or apply class sampling to the $p = 500$ design.

Sign disagreement rises sharply with the ambient number of columns. At $p = 500$, 32.88 percent of observations have opposite signs with both estimates at least 0.10 in magnitude.

Changing only the forest seed produces 0.09 percent sign disagreement. Class sampling removes the difference exactly on the pooled training and evaluation array.

For the $N = 30,000$ evaluation observations $X_1^e, \dots, X_N^e$, define CATE mean squared error (MSE) MSE_j and treatment-rule regret Reg_j as

$$\text{MSE}_j = \frac{1}{N} \sum_{i=1}^N \{\hat{\tau}_j(X_i^e) - \tau(X_i^e)\}^2, \quad \text{Reg}_j = \frac{1}{N} \sum_{i=1}^N |\tau(X_i^e)| \mathbf{1}\{\hat{\pi}_j(X_i^e) \neq \pi^*(X_i^e)\},$$

where j indexes the fitted forest, $\hat{\pi}_j(x) = \mathbf{1}\{\hat{\tau}_j(x) > 0\}$ is its fitted treatment rule, and $\pi^*(x) = \mathbf{1}\{\tau(x) > 0\}$ is the optimal treatment rule. Table 2 reports both losses, including results for class sampling, which was evaluated at $p = 500$. The class-sampled fits use the same oracle nuisance functions, certify weak-order classes on the pooled training and evaluation covariates, and apply the default candidate budget to the number of classes.

Table 2: CATE accuracy and treatment-rule regret					
p	Forest	MSE		Regret	
40	Recodings of X_1	0.0361	(0.0012)	0.0127	(0.0009)
	Recodings of X_2	0.0296	(0.0010)	0.0058	(0.0007)
100	Recodings of X_1	0.0747	(0.0018)	0.0222	(0.0006)
	Recodings of X_2	0.0615	(0.0018)	0.0093	(0.0015)
250	Recodings of X_1	0.1450	(0.0019)	0.0257	(0.0001)
	Recodings of X_2	0.1511	(0.0023)	0.0438	(0.0026)
500	Recodings of X_1	0.1962	(0.0015)	0.0262	(0.0001)
	Recodings of X_2	0.2346	(0.0020)	0.0647	(0.0016)
	Class sampling	0.3653	(0.0021)	0.0047	(0.0004)

Notes: Means across the same 100 replications as the representation comparisons in Table 1. Parentheses give Monte Carlo standard errors. Both class-sampled representations have exactly identical predictions in every replication. Regret is in outcome units per person.

At $p = 500$, class sampling has mean squared error 0.3653, compared with 0.1962 and 0.2346 when recodings of X_1 and X_2 are retained. Its mean regret is 0.0047, compared with 0.0262 and 0.0647. Class sampling improves treatment decisions in this design while increasing mean squared CATE error.

5.2 Empirical Replication: Job Training and Earnings

I replicate specifications 1 and 2 of Chernozhukov et al. [2022] for the National Supported Work (NSW) experimental comparison and the Panel Study of Income Dynamics (PSID) and Current Population Survey (CPS) comparison groups. Across these six cases, the re-

constructed full-dictionary ATETs are within 0.10 published standard errors of the reported estimates.

The outcome-regression dictionaries include age, education, 1974 earnings, and 1975 earnings together with their squares. Since their observed supports are nonnegative, each raw-square pair offers the same CART threshold partitions. I certify classes on the pooled observed and treatment-counterfactual rows, including both treatment states for every covariate vector used in prediction. The reduction retains one original column per class and removes four or five of 24 to 30 columns. It changes only the random-forest outcome regression, leaving the samples, folds, Riesz/Lasso stage, and orthogonal score fixed.

Table 3 compares the full and reduced dictionaries using original column values. One reduction recomputes the `randomForest` regression default, one third of the supplied column count rounded down, after removing columns. The other holds its original numerical value fixed. All fits use 1,000 trees and the same seed. The final column benchmarks representation changes against four alternative seeds of the full dictionary.

Table 3: ATET estimates with full and reduced forest dictionaries			
Sample / specification	Full dictionary	Reduced, default	Reduced, fixed
		mtry	mtry
NSW / 1	3,013.6 (1,277.9)	3,127.8 (1,322.0)	3,043.6 (1,304.9)
NSW / 2	3,207.0 (1,364.9)	3,191.1 (1,359.7)	3,038.1 (1,264.3)
PSID / 1	1,469.0 (975.6)	1,555.0 (924.2)	1,645.8 (941.6)
PSID / 2	1,381.1 (949.0)	1,362.6 (921.7)	1,345.0 (931.0)
CPS / 1	1,636.2 (616.3)	1,483.6 (599.7)	1,621.5 (607.5)
CPS / 2	1,566.2 (618.3)	1,505.1 (605.7)	1,570.0 (608.5)
	Change from full dictionary		Seed variation
	Default mtry	Fixed mtry	Mean absolute (max.)
NSW / 1	+114.3	+30.0	43.8 (86.4)
NSW / 2	-15.9	-169.0	176.1 (234.7)
PSID / 1	+86.0	+176.8	36.4 (61.4)
PSID / 2	-18.5	-36.1	46.1 (64.8)
CPS / 1	-152.6	-14.7	22.3 (30.4)
CPS / 2	-61.1	+3.8	10.1 (20.0)

Notes: All amounts are in dollars. Upper-panel parentheses give ATET standard errors. Full-dictionary entries are reconstructed estimates. Each fit uses 1,000 trees and fixed five-fold cross-fitting. Lower-panel changes pair seed 1 before and after reduction. Seed variation compares the full-dictionary estimate at seed 1 with seeds 2–5. Specifications 1 and 2 have 24 and 30 original columns. Reduction leaves 19 and 25 columns in NSW and CPS, and 20 and 26 in PSID. Recomputed `mtry` is 6 and 8, versus 8 and 10 when held fixed.

The ATET changes range from 4 to 177 dollars in absolute value. All estimates remain positive, with no changes in statistical significance at the 5 percent level. Representation changes do not uniformly exceed seed variation in the aggregate ATET. For example, in NSW

specification 2, the default-budget reduction changes the estimate by 16 dollars, compared with a mean absolute seed change of 176 dollars.

For the cross-fitted outcome regression $\hat{m}(w, x)$, the sign of $\hat{m}(1, x) - \hat{m}(0, x)$ changes for 2.3–11.5 percent of all observations and 3.2–11.4 percent of treated observations. The corresponding seed benchmarks range from 1.3 to 4.6 percent and from 1.8 to 3.5 percent.

Native midpoint cutoffs can route evaluation observations differently after a nonlinear recoding. Table 4 removes this difference by assigning identical dense ranks to equivalent columns on the pooled observed and counterfactual array. Both dictionaries use these same ranks, the original numerical `mtry`, folds, seed, and Riesz estimates. Only the forest inputs are recoded. The orthogonal score retains its original dictionary.

Table 4: ATET comparisons with identical canonical rank coding

Sample / specification	Full	Reduced	Change	Sign changes (%)
NSW / 1	3,142.3 (1,340.2)	3,066.5 (1,290.1)	-75.8	3.36
NSW / 2	3,221.9 (1,410.8)	3,289.5 (1,413.2)	+67.6	2.80
PSID / 1	1,506.4 (971.5)	1,646.2 (941.2)	+139.8	2.52
PSID / 2	1,385.5 (949.7)	1,336.4 (934.9)	-49.1	2.63
CPS / 1	1,677.7 (616.3)	1,597.1 (607.0)	-80.6	6.83
CPS / 2	1,551.6 (618.1)	1,522.5 (609.7)	-29.2	8.34

Notes: ATETs and changes are in dollars, with standard errors in parentheses. Changes are reduced minus full. Both dictionaries use 1,000 trees, seed 1, the same five folds, and `mtry` equal to 8 or 10. Sign changes concern the cross-fitted outcome-regression contrast over all observations.

With canonical ranks, absolute ATET changes range from 29 to 140 dollars, and 2.5–8.3 percent of individual contrasts change sign. The ATET estimates remain positive, with no changes in statistical significance at the 5 percent level. Thus nonlinear midpoint routing does not account for all of the first-stage sensitivity. Column-dependent tie handling and pseudorandom draws can still differ. Neither empirical comparison measures accuracy or regret, since the true conditional effects are unknown.

Code and outputs are available in the [replication repository](#). From its root, `python3 scripts/verify_replication.py` reconstructs the tables. The package documents full re-fits, recodings, retained columns, software versions, folds, and seeds.

6 Conclusion

Feature subsampling makes column allocation part of a forest’s specification. Under the conditions of Theorem 1, reallocating redundant columns can change the limiting CATE and treatment rule without changing covariate information. In the Monte Carlo, class sam-

pling lowers treatment-rule regret but increases CATE mean squared error. The replication illustrates that aggregate ATET comparisons alone can miss changes in individual first-stage contrasts. Sampling certified split classes removes the extra candidate weight assigned to redundant columns. Under common class-level randomness and canonical routing, equivalent representations give identical predictions on the certified array.

A Proofs

Proof of Theorem 1. All probability bounds for the endpoints are uniform over $a \in \mathcal{C}_a$. Write $\mathcal{D}_n$ for the training data and the fitted nuisance functions, including their training randomization. Uniform boundedness, the fixed two-dimensional rectangle class, and relative Vapnik–Chervonenkis (VC) Bernstein bounds give

$$r_n = \sqrt{\frac{\log(B_n s_n)}{k_n}} + \frac{\log(B_n s_n)}{k_n} + \|\widehat{m} - m_a\|_\infty = o_p(1).$$

These bounds hold simultaneously for the root splitting and estimation samples and all their rectangular restrictions. They therefore cover both ordinary and failure trees, even though the latter use conditional transition laws. Increasing recodings do not enlarge this rectangle class. The score identity and its perturbation bound are

$$\mathbb{E}(\Gamma_i \mid X_i = x) = \tau_a(x), \quad \max_i |\widehat{\Gamma}_i - \Gamma_i| \leq 2\|\widehat{m} - m_a\|_\infty.$$

Combining the concentration bounds with (5) yields

$$\sup_{C,j,t} \left| \widehat{\mathcal{Q}}_C^L(j, t) - \mathcal{Q}_C(j, t) \right| = o_p(1), \quad (14)$$

where $\mathcal{Q}_C$ is the population score gain.

First consider a node before coordinate j has been used. Its range in that coordinate is still $[0, 1]$, and independence removes the other component from the split contrast. Direct calculation gives

$$\mathcal{Q}_C(j, t) = \begin{cases} a_j^2 t / (1 - t), & 0 < t \leq 1/2, \\ a_j^2 (1 - t) / t, & 1/2 \leq t < 1. \end{cases} \quad (15)$$

The unique maximizer is $1/2$, with gain a_j^2 . Positivity of $\underline{a}$ and the coefficient gap imply $a_2^2 - a_1^2 \geq 2\underline{a}\Delta_a > 0$. At nodes with at least $C_0 k_n$ splitting observations, cuts approaching $1/2$ are size-admissible and balanced on the concentration event. Maximization and (14) thus imply

$$\max_{j,F} |\widehat{t}_{j,F} - 1/2| = o_p(1), \quad (16)$$

where F ranges over first-use nodes above this size band in the ordinary and failure trees.

The continuation condition now supplies the depth bound. It is used here in addition to

balance. If $n_d(x)$ is the splitting-sample size on a path, then, whenever a split is made,

$$\beta n_d(x) \leq n_{d+1}(x) \leq (1 - \beta)n_d(x).$$

The root size is proportional to s_n , no path stops above $C_0 k_n$, and every child has at least k_n observations. Consequently there are constants $0 < c_\beta < C_\beta < \infty$ and an integer K , independent of n, a, h , such that, with probability tending to one,

$$c_\beta H_n \leq J_{h,n}(X), J_{h,n}^0(X) \leq C_\beta H_n. \quad (17)$$

Here $J_{h,n}$ is the number of candidate draws on an ordinary path, and $J_{h,n}^0$ is defined when $\omega_n < 1$. At most K draws occur after the path enters the band $n_d < C_0 k_n$, including a final unsuccessful search. Reached leaves have splitting sizes between k_n and $C_0 k_n$. Relative concentration gives their population masses of order k_n/s_n and honest estimation counts of order k_n . In particular,

$$\max_{b,L} \left| \frac{1}{N_{b,n}^{\text{est}}(L)} \sum_{i \in I_{b,n}^{\text{est}}: X_i \in L} \hat{\Gamma}_i - \mathbb{E}\{\tau_a(X) \mid X \in L\} \right| = o_p(1). \quad (18)$$

The maximum includes both representations. The counts are positive on this event.

We next control the approximation after first use. Choose a deterministic $\delta_n \downarrow 0$ so that (16) is at most δ_n with probability tending to one. For each coordinate, its first-use nodes above the terminal band form an antichain. Label only their children by the sign suggested by the side of the cut, and retain this label in their descendants. The total population mass receiving a wrong label is

$$\sum_F P_X(F) |\hat{t}_{j,F} - 1/2| \leq \delta_n,$$

where the sum is over these above-band first-use nodes. For a descendant C , let $q(C)$ be the conditional fraction with the wrong label. Then

$$\sup_t \mathcal{Q}_C(j, t) \leq 4\bar{a}^2 q(C), \quad \int_C |\mathbb{E}(s_j \mid C) - s_j(x)| dP_X(x) \leq 4P_X(C)q(C).$$

These bounds also apply after further splits. For $u_n = \sqrt{\delta_n}$, select the first descendant on each path for which $q(C) > u_n$. Those descendants are disjoint, so their combined mass is at most δ_n/u_n . Outside this vanishing set, a previously used coordinate has gain at most $4\bar{a}^2 u_n$ at every subsequent node. An offered unused coordinate has gain at least $\underline{a}^2 - o_p(1)$ above the terminal band and therefore wins. Only the root can involve competition between

two unused coordinates.

The repeated coordinate is offered with probability at least $1/2$ at each draw. Together with (17), the preceding ordering argument bounds its probability of remaining unused above the terminal band by

$$2^{-\lfloor c_\beta H_n \rfloor + K + 1} + o(1) = o(1).$$

For the singleton, the probability of first use only in that band is also $o(1)$. If $\omega_n \leq H_n^{-1/2}$, it is at most $K\omega_n$. Otherwise, avoiding an earlier use has probability at most $\exp\{-\omega_n(\lfloor c_\beta H_n \rfloor - K - 1)\} + o(1)$. The error terms here include the exceptional target mass just bounded.

Let $D_{h,n}(x)$ indicate first use of the singleton in a generic tree, and let $\check{T}_{h,n}^{\text{tree}}$ be that tree's honest score prediction. If $D_{h,n}(x) = 0$, its leaf has no restriction in the singleton coordinate, whose conditional mean is exactly zero. For a first use above the terminal band, localization and the wrong-label bound control the integrated approximation error. A late first use adds at most twice its probability to the integrated error for the sign component, which vanishes by the preceding bound. The same bound applies to the repeated component. Thus

$$\mathbb{E}_\Theta \left[\left\| \check{T}_{h,n}^{\text{tree}} - D_{h,n}\tau_a - (1 - D_{h,n})a_h s_h \right\|_{L^1(P_X)} \middle| \mathcal{D}_n \right] = o_p(1). \quad (19)$$

This is an approximation after first use, not an assertion that every leaf is exactly pure.

Suppose now that $\omega_n < 1$. Conditional on $Q_{M_n} = q$, the singleton is included with probability q/M_n , so its unconditional inclusion probability at a node is ω_n . On the locally conditioned failure tree, write $\rho_{h,n,t}^0(x)$ for the probability that the ordinary transition at the current node chooses the singleton as its splitting coordinate. It satisfies $0 \leq \rho_{h,n,t}^0 \leq \omega_n$. After the root and above the terminal band, this probability equals ω_n outside the exceptional paths. With $\Lambda_{h,n}^0 = \sum_{t \leq J_{h,n}^0} \rho_{h,n,t}^0$, we obtain

$$|\Lambda_{h,n}^0 - J_{h,n}^0 \omega_n| \leq (K + 1)\omega_n \quad (20)$$

outside a set of joint training, failure-tree, and test-point probability tending to zero.

Node locality ensures that conditioning transitions on other branches leaves the transition kernels on the target path unchanged. At each stage, factor the probability of avoiding a singleton split from the transition kernel conditional on that event. Multiplication along the path and integration over the conditional kernels give the exact identity

$$u_{h,n}(x) := \mathbb{P}_\Theta \{D_{h,n}(x) = 0 \mid \mathcal{D}_n\} = \mathbb{E}_{\Theta^0} \left[\prod_{t=1}^{J_{h,n}^0(x)} (1 - \rho_{h,n,t}^0(x)) \middle| \mathcal{D}_n \right]. \quad (21)$$

The kernels are well-defined because their conditioning probabilities are at least $1 - \omega_n > 0$. This identity concerns no first use throughout the path, exactly the event defining $D_{h,n} = 0$.

If $\mathcal{I}_n \rightarrow 0$, then $\omega_n \rightarrow 0$ and $\Lambda_{h,n}^0 \rightarrow_p 0$ by (17) and (20). The product converges to one since its deficit is at most $\Lambda_{h,n}^0$. If $\mathcal{I}_n \rightarrow \infty$ and $\omega_n < 1$, the cumulative hazard diverges and

$$\prod_t (1 - \rho_{h,n,t}^0) \leq e^{-\Lambda_{h,n}^0} \rightarrow_p 0.$$

For indices with $\omega_n = 1$, both coordinates are always offered. The stronger one wins at the root and the other at the next interior node, outside the same exceptional set. Thus recovery also holds on these indices, without a failure-tree conditioning on a null event.

For the intermediate case, fix $a \in \mathcal{C}_a$. The condition $\mathcal{I}_n \rightarrow \kappa \in (0, \infty)$ implies $\omega_n \rightarrow 0$. Equations (7) and (20) yield $\Lambda_{h,n}^0 \rightarrow_p \kappa d_h(a)$. Moreover,

$$0 \leq -\log \prod_t (1 - \rho_{h,n,t}^0) - \Lambda_{h,n}^0 \leq \frac{\omega_n \Lambda_{h,n}^0}{1 - \omega_n} \rightarrow_p 0.$$

The product therefore converges to $e^{-\kappa d_h(a)}$. All products are bounded. Joint convergence with an independent test point implies convergence of $u_{h,n}$ in $L^1(P_X)$, in probability over $\mathcal{D}_n$, to one, zero, or $e^{-\kappa d_h(a)}$ in the respective regimes.

Finally, conditional on $\mathcal{D}_n$, the independent partition of the B_n trees into groups of bounded size, together with bounded score predictions, gives

$$\|\check{T}_{h,n} - \mathbb{E}_{\Theta}(\check{T}_{h,n}^{\text{tree}} \mid \mathcal{D}_n)\|_{L^1(P_X)} = O_p(B_n^{-1/2}) = o_p(1).$$

Combine this with (19), the three limits for $u_{h,n}$, and the prediction transfer in (5). This proves (8), (9), and (10).

A strictly increasing recoding maps a threshold t to $\varphi(t)$ and preserves $\varphi(x) \leq \varphi(t)$ exactly when $x \leq t$. All cells and concentration bounds above are expressed in the original coordinates. Their constants therefore do not depend on the number or form of the increasing recodings. $\square$

Proof of Corollary 1. In the selection regime, Markov's inequality and (8) give, for $g = 1, 2$,

$$\mathbb{P}_X\{\text{sgn}(\hat{T}_{g,n}) \neq s_g\} \leq \frac{2}{\underline{a}} \|\hat{T}_{g,n} - a_g s_g\|_{L^1(P_X)} = o_p(1).$$

The independent signs s_1 and s_2 disagree with probability one half. On that event, $|\tau_a| = a_2 - a_1 \geq \Delta_a$. A second application of Markov's inequality shows that both fitted magnitudes exceed $\underline{a}/2$ outside a set of P_X -mass $o_p(1)$. This proves (11) and the stated separation from

zero.

For any rule π ,

$$\text{Reg}_a(\pi) = \mathbb{E}[|\tau_a(X)|\mathbf{1}\{\pi(X) \neq \pi_a^*(X)\}].$$

Since $a_2 > a_1$, $\pi_a^* = \mathbf{1}\{s_2 = 1\}$, so the limit of $\widehat{\pi}_{2,n}$ has zero regret. For the other representation,

$$\text{Reg}_a(\mathbf{1}\{s_1 = 1\}) = \mathbb{E}[\tau_a\{\mathbf{1}\{s_2 = 1\} - \mathbf{1}\{s_1 = 1\}\}] = \frac{a_2 - a_1}{2}.$$

Boundedness and classification convergence transfer these values to the estimated rules, proving (12). $\square$

Proof of Corollary 2. The capped-Poisson rule satisfies

$$\mathbb{E}Q_M = \xi_M + e^{-\xi_M} - \mathbb{E}(N_M - M)_+.$$

For fixed $\xi_M = \xi$, the final term vanishes as $M \rightarrow \infty$, giving $\omega_n \sim (\xi + e^{-\xi})/M_n$. If $\xi_M \rightarrow \infty$ and $\xi_M/M \rightarrow 0$, Poisson concentration gives $\mathbb{E}Q_M/\xi_M \rightarrow 1$. Under the package default, $\xi_M/\sqrt{M} \rightarrow 1$, so $\omega_n \sim M_n^{-1/2}$. The two boundaries follow by substituting these expressions into $\mathcal{I}_n = H_n\omega_n$. For fixed finite M , $\omega_n > 0$ is constant and $H_n \rightarrow \infty$. $\square$

Proof of Proposition 1. Couple the two representations with the common class-indexed randomness in the proposition. At a common node they draw the same classes, evaluate the same certified actions, and apply the same scores and tie order. Equal certificates give identical unordered children, while canonical orientation sends every point in $\mathcal{X}_{\text{cert}}$ to the same child. Induction gives identical tree partitions on the certified array. Leaf observations and forest weights therefore coincide. Nuisance fits coincide because they are held fixed or computed from the same canonical matrix with common randomness. The same inputs, settings, and randomization then give identical stopping decisions, predictions, variance calculations, rankings, and tie-broken selected sets. $\square$

References

- Susan Athey, Julie Tibshirani, and Stefan Wager. Generalized random forests. *Annals of Statistics*, 47(2):1148–1178, 2019. doi: 10.1214/18-AOS1709.
- Zheng Chen and Weixiong Zhang. Integrative analysis using module-guided Random Forests reveals correlated genetic factors related to mouse weight. *PLoS Computational Biology*, 9(3):e1002956, 2013. doi: 10.1371/journal.pcbi.1002956.

- Victor Chernozhukov, Whitney K. Newey, and Rahul Singh. Automatic debiased machine learning of causal and structural effects. *Econometrica*, 90(3):967–1027, 2022. doi: 10.3982/ECTA18515.
- Chien-Ming Chi, Patrick Vossler, Yingying Fan, and Jinchi Lv. Asymptotic properties of high-dimensional random forests. *Annals of Statistics*, 50(6):3415–3438, 2022. doi: 10.1214/22-AOS2234.
- Christian Colot, Philippe Baecke, and Isabelle Linden. Leveraging fine-grained mobile data for churn detection through essence random forest. *Journal of Big Data*, 8(1):63, 2021. doi: 10.1186/s40537-021-00451-9.
- GRF Development Team. The GRF algorithm. Algorithm reference for the `grf` package, 2026. URL <https://grf-labs.github.io/grf/REFERENCE.html>. Accessed August 6, 2026.
- Jason M. Klusowski and Peter M. Tian. Large scale prediction with decision trees. *Journal of the American Statistical Association*, 119(545):525–537, 2024. doi: 10.1080/01621459.2022.2126782.
- Gilles Louppe. *Understanding Random Forests: From Theory to Practice*. PhD thesis, University of Liège, 2014. URL <https://orbi.uliege.be/handle/2268/170309>.
- Tianxing Mei, Yingying Fan, and Jinchi Lv. Exogenous randomness empowering random forests. *Journal of Machine Learning Research*, 27(98):1–95, 2026. URL <https://jmlr.org/papers/v27/25-0247.html>.
- Stefan Wager and Susan Athey. Estimation and inference of heterogeneous treatment effects using random forests. *Journal of the American Statistical Association*, 113(523):1228–1242, 2018. doi: 10.1080/01621459.2017.1319839.